

Review Questions

2. a) { EventID → EventJournal (all attributes) }

b) { JourISBN → JournalName, JournalName → JourISBN, JourISBN → Publisher }

c) The table is not in 3NF. The functional dependency JourISBN → JournalName is a non-trivial functional dependency where JourISBN is not a superkey of EventJournal, and JournalName → JourISBN is not contained in a candidate key of EventJournal. We can use the Compute3NF algorithm to create a 3NF decomposition of EventJournal.

We compute a minimal cover:

F = {EventID → EventName, EventWebLink, JournalName, Publisher, JourISBN;
JourISBN → JournalName;
JournalName → JourISBN;
JourISBN → Publisher }

Fc = {EventID → EventName, EventWebLink, Publisher, JourISBN;
JourISBN → JournalName;
JournalName → JourISBN }

The candidate keys are EventID. We decompose the table into

{(EventID, EventName, EventWebLink, Publisher, JourISBN), (JourISBN, JournalName)}

3. This is not a valid approach, since every non-primary key needs to be fully dependent on the primary key, and O and X are not fully dependent on PK, Y.

4. a) ContractID → CustomerID, ServiceAddress, ServiceType, StartDate
CustomerID → ServiceAddress

b) We use the primary key ContractID. This is a superkey of R, and every non-key is fully dependent on it.

c) I didn't list these in part (a) bc I only listed the non-redundant ones so here ya go:
ContractID → CustomerID
ContractID → CustomerID, ServiceAddress
ContractID → CustomerID, ServiceAddress, ServiceType
ContractID → ServiceAddress
...
I'm sure you get the idea

5. a) Each record takes 15+2+10+10 = 37 bytes. (What does fields starting at a byte that is a multiple of 8 have to do with this?)

b) floor(512/37) = 13. Each block can store 13 records. We need floor(5000/13) = 385 blocks.

c) In the worst-case, we search all the blocks. This would be 385 block accesses.

d) In the worst-case, we search all the blocks. This would be 385 block accesses.

8. Conflict-serializable schedule:
r4[b], r3[a], r2[b], r1[a], w2[b]
r1[d], w1[d],
r2[d], w2[d],
w3[a], r3[c], w3[c],
w4[b], r4[a], w4[a]

Equivalent serial schedule:
r1[a], r1[d], w1[d],
r2[b], w2[b], r2[d], w2[d],
r3[a], w3[a], r3[c], w3[c],
r4[b], w4[b], r4[a], w4[a]

6. Result (same as Temp):

Rating	AvgAge
1	33
3	38.17
7	40
8	40.5
9	35
10	25.5

8. a) T2read(Z), T3write(Z)
T2read(Y), T3write(Y)

T2write(Y), T3read(Y)
T2write(Y), T3write(Y)
T2write(Y), T1read(Y)
T2write(Y), T1write(Y)

T2read(Y), T1write(Y)
T3read(Y), T1write(Y)

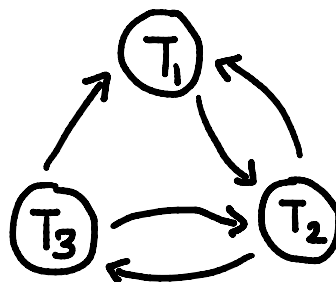
T1read(X), T2write(X)

T1write(X), T2read(X)
T1write(X), T2write(X)

T3write(Y), T1read(Y)
T3write(Y), T1write(Y)

T3write(Z), T2read(Z)

b) This schedule is not conflict-serializable, since there are cycles in its serialization graph.



d) T3 reads item Y after it is written by T2 before T2 has committed (nevermind the fact that there are no commit operations at all), so it is not a recoverable schedule.

1. b) Email -> Address
CardNumber, ExpiryDate -> Email, Address
CardNumber -> ExpiryDate
Name -> Price, Duration
Email, Name -> RentDuration, Discount, StartTime

c) A primary key could be Email, CardNumber, Name. It is a candidate key for this relation, is unique, and cannot be reduced.

d) This is not compliant with 2NF. The non-prime attribute ExpiryDate is dependent on just CardNumber, which is a subset of the primary key. We can decompose the relation as follows:

CardNumber, ExpiryDate
Email, CardNumber
Email, Address
Name, Price, Duration
Email, Name, RentDuration, Discount, StartTime

2. a) 3NF ???

b) False. It is possible to have every non-prime attribute fully dependent on each candidate key. For example, we can have the relation (A, B, C) where A and B are candidate keys. So we have $A \rightarrow BC$, $B \rightarrow AC$. C is the only non-prime attribute. It is fully dependent on A, and fully dependent on B. Hence this is in 2NF.

c) We have $P \rightarrow CX$, $C \rightarrow PX$. There are no transitive dependencies.

3. a) Yes. We can have $X \rightarrow C1$. C1 is in a candidate key, so this does not violate 3NF. But X is not in {P, C1, C2}.

8. a)

```
select Event.EventName, count(TopicID)
from Event
join Covers on Event.EventID = Covers.EventID
group by EventID
```

c)

```
select Event.EventName, count(TopicID)
from Event
join Covers on Event.EventID = Covers.EventID
where substr(Event.EventName, 0, 1) == 'A'
group by EventID
```

d)

```
select Event.EventName, count(TopicID)
from Event
join Covers on Event.EventID = Covers.EventID
group by EventID
having count(TopicID) < 5
```

g)
 select EventName
 from Event
 where EventID not in (
 select EventID
 from covers
)

l. a)
 (slide 17-18)
 select SalaryGrades.HighestSal - Employees.salary
 from SalaryGrades
 join Employees on Employees.salary >= SalaryGrades.LowestSal and Employees.salary <= SalaryGrades.HighestSal
 where Employees.EmpName = 'Smith'

b)
 select EmpName
 from Employees
 where salary >= (
 select LowestSal
 from SalaryGrades
 where grade = 'G1'
)
 and salary <= (
 select HighestSal
 from SalaryGrades
 where grade = 'G1'
)

7.
 (slide 19)
 i. Superkey
 ii. Both

8.
 someOtherColumn could have a NULL value (I am pretty sure this does not throw an error though, might just have unexpected behavior since NULL > 5 will eventually evaluate to false?)

9.
 Yes.

4. b)
 select model
 from (
 (select model, price from PC)
 union
 (select model, price from Printer)
)
 where price >= all (
 (select model, price from PC)
 union
 (select model, price from Printer)
)