

ECE459 Final W2011

Question 1

- a) Not covered.
- b) Throughput
- c) Memory accesses
- d) CUDA equivalent: work items
- e) Problem size
- f) Array aliasing
- g) Minimal performance impact when not being actively used
- h) Safety and reliability
- i) Stores
- j) Not covered.
- k) Not covered.
- l) Compile sources with prof-gen option, run instrumented executable, compile with prof-use option
- m) Not covered.
- n) Mutual exclusion
- o) Not covered. Register renaming
- p) Profile
- q) Parallelization
- r) Acquire the same lock multiple times without deadlock
- s) Client-server
- t) Kernel

Question 2

Not covered.

Question 3

Not covered.

Question 4

Since the architecture is not sequentially consistent, operations may be reordered.

Final values	Execution order
r1=0, r2=1, x=1, y=1	x=1; r1=y; y=x; r2=x (T1; T2) (can swap first 2 ops for same result)
r1=0, r2=0, x=1, y=0	y=x; r2=x; x=1; r1=y (T2; T1) (can swap last 2 or first 2 ops for same result)
r1=1, r2=1, x=1, y=1	x=1; y=x; r1=y; r2=x (can swap last 2 ops for same result)
r1=0, r2=1, x=1, y=0	y=x; x=1; r1=y; r2=x (can swap last 2 ops for same result)
r1=0, r2=0, x=1, y=1	r1=y; r2=x; x=1; y=x (can swap first 2 ops for same result)

Question 5

One compiler optimization could be loop unrolling: we could process 4 elements at a time in the inner loop by incrementing *i* by 4 each time instead and updating the summation accordingly. This requires the condition that there are no loop-carried dependencies, which there aren't.

Another compiler optimization could be function inlining: we could inline the max function at line 17.

Question 6

Written in CUDA instead of OpenCL. The atomic max function is used to ensure atomic writes to the result.

```
__global__ void kernel1 (double *u, double *w, int M, int N) {  
    int i = blockIdx.x*blockDim.x + threadIdx.x + 1;  
    int j = blockIdx.y*blockDim.y + threadIdx.y + 1;  
    w[ i *N+j ] = ( u [ (i - 1)*N +j ] + u [( i + 1)*N+ j ] + u [ i *N + j -1] + u [ i *N + j  
+1] ) / 4 . 0;  
}  
  
__global__ void kernel2(double *u, double *w, int M, int N, double *diff) {  
    int i = blockIdx.x*blockDim.x + threadIdx.x + 1;  
    int j = blockIdx.y*blockDim.y + threadIdx.y + 1;  
    double diffNew = fabs(w[i*N+j] - u[i*N+j]);  
    atomicMax(diff, diffNew);  
}
```